

Supervised Classification Models for Text Analysis

Bamberg Summer Institute in Computational Social Science


Carsten Schwemmer, University of Bamberg

2019-08-01


Supervised text classification

Supervised vs unsupervised

Supervised Learning



Unsupervised Learning



A simple model: Naive-Bayes Classifier

- “all models are wrong but some models are useful.”
- *naive* assumption: features are independent of each other

$$P(\text{label} \mid \text{feature}) = \frac{P(\text{feature} \mid \text{label}) \cdot P(\text{label})}{P(\text{feature})}$$

Reload our DonorsChoose data

```
library(tidyverse)
library(quantda)
library(caret)
load('data/dfm_donor.Rdata')
dim(dfm_donor)
```

```
## [1] 10000 6751
```

We want to predict whether a donation request received funding:

```
docvars(dfm_donor) %>% count(funded)
```

```
## # A tibble: 2 x 2
##   funded      n
##   <dbl> <int>
## 1      0 2409
## 2      1 7591
```

- in order to train a model, we need to split our data into training and test sets
- the classifier will learn from the training set. The testset is used to evaluate its performance on unseen data
- there is no optimal solution for the proportions to split train and test data (see bias-variance tradeoff).

Training a model

```
set.seed(1337) # for replication

train <- dfm_sample(dfm_donor, size = 8000) # 80% of the data
test <- dfm_donor[!docnames(dfm_donor) %in% docnames(train), ]

# run the model
donor_nb <- textmodel_nb(x = train, y = docvars(train)$funded)
```

Making sense of the model - important features

We create a function to find the most important terms for correctly predicting whether a donation request received funding:

```
imp_features <- function(nb_model, n = 10) {  
  # PcGw = probability of class given the word  
  features <- t(nb_model$PcGw) %>% # transpose  
    as.data.frame() %>% rownames_to_column('feature') %>%  
    gather('label', 'prob_class' , -feature) %>% # tidy data  
    group_by(label) %>% top_n(n, prob_class) %>% # max probability  
    slice(1:n) %>%  
    ungroup() %>% arrange(label, desc(prob_class)) # sort  
  return(features)  
}
```


Making sense of the model - important features

```
imp_features(donor_nb, n = 5)
```

```
## # A tibble: 10 x 3
```

##	feature	label	prob_class
##	<chr>	<chr>	<dbl>
##	1 intuitive	0	0.903
##	2 rated	0	0.879
##	3 france	0	0.875
##	4 associate	0	0.875
##	5 parcc	0	0.875
##	6 harvey	1	0.965
##	7 failing	1	0.896
##	8 teammates	1	0.893
##	9 renewable	1	0.893
##	10 arrived	1	0.893

Confusion matrix & accuracy

	Reference	
Predicted	<i>TRUE</i>	<i>FALSE</i>
<i>TRUE</i>	True Positives	False Positives
<i>FALSE</i>	False Negatives	True Negatives

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + FN + TN}$$

- accuracy is often not a good measure of performance, especially for imbalanced classes.
- multiple alternatives are available, e.g. balanced accuracy

$$\text{Sensitivity} = \frac{TP}{TP + FN}; \text{Specificity} = \frac{TN}{TN + FP}$$

$$\text{Balanced Accuracy} = \frac{\text{Sensitivity} + \text{Specificity}}{2}$$

Evaluating model performance

```
pred_labels <- predict(donor_nb, newdata = test) %>% as.factor()
true_labels <- docvars(test)$funded %>% as.factor()
stats <- confusionMatrix(pred_labels, true_labels)
stats$table # confusion matrix
```

```
##           Reference
## Prediction    0    1
##           0  227  502
##           1  244 1027
```

```
stats$byClass # performance metrics
```

```
##           Sensitivity           Specificity      Pos Pred Value
##           0.4819533           0.6716808           0.3113855
##           Neg Pred Value           Precision           Recall
##           0.8080252           0.3113855           0.4819533
##           F1           Prevalence      Detection Rate
##           0.3783333           0.2355000           0.1135000
## Detection Prevalence      Balanced Accuracy
##           0.3645000           0.5768171
```

- repeat the model training on several train/test splits
- assess performance across all runs
- we'll use a simple k-fold variant, where k equals the number of splits

Running validation over 10 splits

```
cross_val(input_dfm = dfm_donor, labels = 'funded' , train_size = 9000,  
          nr_runs = 10, what = 'Balanced Accuracy', pos_class = '1')
```

```
## [1] "Balanced Accuracy run 1 : 0.589"  
## [1] "Balanced Accuracy run 2 : 0.571"  
## [1] "Balanced Accuracy run 3 : 0.585"  
## [1] "Balanced Accuracy run 4 : 0.603"  
## [1] "Balanced Accuracy run 5 : 0.573"  
## [1] "Balanced Accuracy run 6 : 0.59"  
## [1] "Balanced Accuracy run 7 : 0.544"  
## [1] "Balanced Accuracy run 8 : 0.571"  
## [1] "Balanced Accuracy run 9 : 0.584"  
## [1] "Balanced Accuracy run 10 : 0.581"  
## [1] "Average Balanced Accuracy : 0.579"
```

What to do when model performance is bad

- adjust feature space
- adjust model parameters
- try different model(s)
- in our case: think about other factors that might be related to the funding success of donation requests

Other approaches we didn't cover

- averaging predictions of multiple models (e.g. support vector machines, random forest models)
- optimizing hyperparameters (e.g. distribution assumptions of our naive-bayes model)
- using features beyond bag-of-words (e.g. word embeddings)

When should you use supervised classification?

- if your goal is the best predictive power, supervised models are reasonable choices
- for many social science applications, supervised models are used to infer labels for a larger dataset from a smaller trainingset, which are then used in downstream tasks
- it might be helpful to compare results with dictionary-based approaches

Questions?